

# Capstone Project-II

CREDIT CARD FRAUD  
DETECTION PROJECT  
REPORT

- By Amit Maurya

# Credit Card Fraud Detection Project Report

## Contents

<b>1</b>	<b>Background</b>	<b>2</b>
1.1	Objectives .....	2
1.2	Structure .....	2
<b>2</b>	<b>Summary</b>	<b>2</b>
2.1	Understanding the Data .....	2
2.2	Preparing the Data .....	6
2.3	Splitting the Data into Train and Test Sets .....	6
2.4	Creating Sampled Sets of Data .....	7
<b>3</b>	<b>Modeling Approach</b>	<b>8</b>
3.1	CART Method .....	9
3.2	Logistic Regression (GLM) Model .....	13
3.3	Random Forest (RF Fit) Model .....	14
3.4	xgboost (XGB Fit) Model .....	15
<b>4</b>	<b>Results</b>	<b>17</b>
<b>5</b>	<b>Conclusion</b>	<b>17</b>

# 1 Background

The USA is the global leader as the most credit card fraud prone country with 38.6 percent of reported fraud losses in 2018. Credit card fraud is the most common and popular kind of identity theft and makes up 35.4% of all identity theft reports. Fraud detection is a challenging problem since fraudulent transactions are rare and they represent a small fraction of transactions but can quickly turn into large sums of money. The good news is that with advances in Machine Learning, system can learn, adapt and uncover emerging patterns for preventing fraud.

## 1.1 Objectives

The objective of this project is to train a Machine Learning algorithm on a dataset made up of credit card transactions in order to successfully predict fraudulent transactions.

## 1.2 Structure

This document is structured as follows:

1. **Introduction/overview:** Introduces the problem and describe the goal of the project.
2. **Summary:** Describes the dataset and variables, explores the data further and prepares data for analysis.
3. **Method/Analysis:** Explains the process and techniques used. Defines the models as we improve upon the model and explains different Machine Learning algorithms and the results for each method.
4. **Results:** Summary of findings presents the modeling results and discusses the model performance.
5. **Conclusion:** Gives a brief summary of the report, its potential impact, its limitations, and future work.

# 2 Summary

## 2.1 Understanding the Data

As a pre-requisite, install all the required packages and import libraries and dependencies for the dataframe. We start by downloading the csv file called creditcard.csv from Kaggle website (Source - <https://www.kaggle.com/mlg-ulb/creditcardfraud>). This dataset was collected and analyzed during a research collaboration of Worldline and the Machine Learning Group (<http://mlg.ulb.ac.be>) of ULB (Université Libre de Bruxelles) on big data mining and fraud detection. The dataset is made up of credit card transactions that occurred in two days in September 2013 by European cardholders.

We read the csv file into a dataframe called 'df' and quickly analyze the structure of the dataframe.

*# We first look at top 6 rows of the dataframe.*

`head(df)`

##	Time	V1	V2	V3	V4	V5	V6
## 1	0	-1.3598071	-0.07278117	2.5363467	1.3781552	-0.33832077	0.46238778
## 2	0	1.1918571	0.26615071	0.1664801	0.4481541	0.06001765	-0.08236081
## 3	1	-1.3583541	-1.34016307	1.7732093	0.3797796	-0.50319813	1.80049938
## 4	1	-0.9662717	-0.18522601	1.7929933	-0.8632913	-0.01030888	1.24720317
## 5	2	-1.1582331	0.87773675	1.5487178	0.4030339	-0.40719338	0.09592146
## 6	2	-0.4259659	0.96052304	1.1411093	-0.1682521	0.42098688	-0.02972755
##		V7	V8	V9	V10	V11	V12
## 1	0	0.23959855	0.09869790	0.3637870	0.09079417	-0.5515995	-0.61780086
## 2	-0	0.07880298	0.08510165	-0.2554251	-0.16697441	1.6127267	1.06523531
## 3	0	0.79146096	0.24767579	-1.5146543	0.20764287	0.6245015	0.06608369

```
## 4  0.23760894  0.37743587 -1.3870241 -0.05495192 -0.2264873  0.17822823
## 5  0.59294075 -0.27053268  0.8177393  0.75307443 -0.8228429  0.53819555
## 6  0.47620095  0.26031433 -0.5686714 -0.37140720  1.3412620  0.35989384
##           V13           V14           V15           V16           V17           V18
## 1 -0.9913898 -0.3111694  1.4681770 -0.4704005  0.20797124  0.02579058
## 2  0.4890950 -0.1437723  0.6355581  0.4639170 -0.11480466 -0.18336127
## 3  0.7172927 -0.1659459  2.3458649 -2.8900832  1.10996938 -0.12135931
## 4  0.5077569 -0.2879237 -0.6314181 -1.0596472 -0.68409279  1.96577500
## 5  1.3458516 -1.1196698  0.1751211 -0.4514492 -0.23703324 -0.03819479
## 6 -0.3580907 -0.1371337  0.5176168  0.4017259 -0.05813282  0.06865315
##           V19           V20           V21           V22           V23           V24
## 1  0.40399296  0.25141210 -0.018306778  0.277837576 -0.11047391  0.06692807
## 2 -0.14578304 -0.06908314 -0.225775248 -0.638671953  0.10128802 -0.33984648
## 3 -2.26185710  0.52497973  0.247998153  0.771679402  0.90941226 -0.68928096
## 4 -1.23262197 -0.20803778 -0.108300452  0.005273597 -0.19032052 -1.17557533
## 5  0.80348692  0.40854236 -0.009430697  0.798278495 -0.13745808  0.14126698
## 6 -0.03319379  0.08496767 -0.208253515 -0.559824796 -0.02639767 -0.37142658
##           V25           V26           V27           V28 Amount Class
## 1  0.1285394 -0.1891148  0.133558377 -0.02105305 149.62      0
## 2  0.1671704  0.1258945 -0.008983099  0.01472417   2.69      0
## 3 -0.3276418 -0.1390966 -0.055352794 -0.05975184 378.66      0
## 4  0.6473760 -0.2219288  0.062722849  0.06145763 123.50      0
## 5 -0.2060096  0.5022922  0.219422230  0.21515315  69.99      0
## 6 -0.2327938  0.1059148  0.253844225  0.08108026   3.67      0
```

*# We then use str to know more about the dataframe and its constituents.*

```
str(df)
```

```
## "data.frame":    284807 obs. of  31 variables:
## $ Time   : num  0 0 1 1 2 2 4 7 7 9 ...
## $ V1     : num -1.36 1.192 -1.358 -0.966 -1.158 ...
## $ V2     : num -0.0728 0.2662 -1.3402 -0.1852 0.8777 ...
## $ V3     : num  2.536 0.166 1.773 1.793 1.549 ...
## $ V4     : num  1.378 0.448 0.38 -0.863 0.403 ...
## $ V5     : num -0.3383 0.06 -0.5032 -0.0103 -0.4072 ...
## $ V6     : num  0.4624 -0.0824 1.8005 1.2472 0.0959 ...
## $ V7     : num  0.2396 -0.0788 0.7915 0.2376 0.5929 ...
## $ V8     : num  0.0987 0.0851 0.2477 0.3774 -0.2705 ...
## $ V9     : num  0.364 -0.255 -1.515 -1.387 0.818 ...
## $ V10    : num  0.0908 -0.167 0.2076 -0.055 0.7531 ...
## $ V11    : num -0.552 1.613 0.625 -0.226 -0.823 ...
## $ V12    : num -0.6178 1.0652 0.0661 0.1782 0.5382 ...
## $ V13    : num -0.991 0.489 0.717 0.508 1.346 ...
## $ V14    : num -0.311 -0.144 -0.166 -0.288 -1.12 ...
## $ V15    : num  1.468 0.636 2.346 -0.631 0.175 ...
## $ V16    : num -0.47 0.464 -2.89 -1.06 -0.451 ...
## $ V17    : num  0.208 -0.115 1.11 -0.684 -0.237 ...
## $ V18    : num  0.0258 -0.1834 -0.1214 1.9658 -0.0382 ...
## $ V19    : num  0.404 -0.146 -2.262 -1.233 0.803 ...
## $ V20    : num  0.2514 -0.0691 0.525 -0.208 0.4085 ...
## $ V21    : num -0.01831 -0.22578 0.248 -0.1083 -0.00943 ...
## $ V22    : num  0.27784 -0.63867 0.77168 0.00527 0.79828 ...
## $ V23    : num -0.11 0.101 0.909 -0.19 -0.137 ...
## $ V24    : num  0.0669 -0.3398 -0.6893 -1.1756 0.1413 ...
## $ V25    : num  0.129 0.167 -0.328 0.647 -0.206 ...
```

```
## $ V26 : num -0.189 0.126 -0.139 -0.222 0.502 ...
## $ V27 : num 0.13356 -0.00898 -0.05535 0.06272 0.21942 ...
## $ V28 : num -0.0211 0.0147 -0.0598 0.0615 0.2152 ...
## $ Amount: num 149.62 2.69 378.66 123.5 69.99 ...
## $ Class : int 0 0 0 0 0 0 0 0 0 0 ...
```

The dataframe has 284,807 transactions with 31 columns (variables). The dataset contains only numerical variables. As stated in the Kaggle website this is a result of a PCA (Principal component analysis) dimensionality reduction to protect sensitive information. Features V1, V2, ... V28 are obtained by PCA except for 'Time', 'Amount', and 'Class'. The variable 'Class' indicates whether a transaction is fraudulent (1 = Fraud) or not (0 = Legal).

summary(df)

```
##      Time      V1      V2      V3
## Min.   : 0      Min.  :-56.40751  Min.   : -72.71573  Min.   : -48.3256
## 1st Qu.: 54202  1st Qu.: -0.92037  1st Qu.: -0.59855  1st Qu.: -0.8904
## Median : 84692  Median : 0.01811  Median : 0.06549  Median : 0.1799
## Mean   : 94814  Mean   : 0.00000  Mean   : 0.00000  Mean   : 0.0000
## 3rd Qu.: 139320 3rd Qu.: 1.31564  3rd Qu.: 0.80372  3rd Qu.: 1.0272
## Max.   : 172792  Max.   : 2.45493  Max.   : 22.05773  Max.   : 9.3826
##      V4      V5      V6      V7
## Min.   : -5.68317  Min.   : -113.74331  Min.   : -26.1605  Min.   : -43.5572
## 1st Qu.: -0.84864  1st Qu.: -0.69160  1st Qu.: -0.7683  1st Qu.: -0.5541
## Median : -0.01985  Median : -0.05434  Median : -0.2742  Median : 0.0401
## Mean   : 0.00000  Mean   : 0.00000  Mean   : 0.0000  Mean   : 0.0000
## 3rd Qu.: 0.74334  3rd Qu.: 0.61193  3rd Qu.: 0.3986  3rd Qu.: 0.5704
## Max.   : 16.87534  Max.   : 34.80167  Max.   : 73.3016  Max.   : 120.5895
##      V8      V9      V10     V11
## Min.   : -73.21672  Min.   : -13.43407  Min.   : -24.58826  Min.   : -4.79747
## 1st Qu.: -0.20863  1st Qu.: -0.64310  1st Qu.: -0.53543  1st Qu.: -0.76249
## Median : 0.02236  Median : -0.05143  Median : -0.09292  Median : -0.03276
## Mean   : 0.00000  Mean   : 0.00000  Mean   : 0.00000  Mean   : 0.00000
## 3rd Qu.: 0.32735  3rd Qu.: 0.59714  3rd Qu.: 0.45392  3rd Qu.: 0.73959
## Max.   : 20.00721  Max.   : 15.59500  Max.   : 23.74514  Max.   : 12.01891
##      V12     V13     V14     V15
## Min.   : -18.6837  Min.   : -5.79188  Min.   : -19.2143  Min.   : -4.49894
## 1st Qu.: -0.4056  1st Qu.: -0.64854  1st Qu.: -0.4256  1st Qu.: -0.58288
## Median : 0.1400  Median : -0.01357  Median : 0.0506  Median : 0.04807
## Mean   : 0.0000  Mean   : 0.00000  Mean   : 0.0000  Mean   : 0.00000
## 3rd Qu.: 0.6182  3rd Qu.: 0.66251  3rd Qu.: 0.4931  3rd Qu.: 0.64882
## Max.   : 7.8484  Max.   : 7.12688  Max.   : 10.5268  Max.   : 8.87774
##      V16     V17     V18
## Min.   : -14.12985  Min.   : -25.16280  Min.   : -9.498746
## 1st Qu.: -0.46804  1st Qu.: -0.48375  1st Qu.: -0.498850
## Median : 0.06641  Median : -0.06568  Median : -0.003636
## Mean   : 0.00000  Mean   : 0.00000  Mean   : 0.000000
## 3rd Qu.: 0.52330  3rd Qu.: 0.39968  3rd Qu.: 0.500807
## Max.   : 17.31511  Max.   : 9.25353  Max.   : 5.041069
##      V19     V20     V21
## Min.   : -7.213527  Min.   : -54.49772  Min.   : -34.83038
## 1st Qu.: -0.456299  1st Qu.: -0.21172  1st Qu.: -0.22839
## Median : 0.003735  Median : -0.06248  Median : -0.02945
## Mean   : 0.000000  Mean   : 0.00000  Mean   : 0.00000
## 3rd Qu.: 0.458949  3rd Qu.: 0.13304  3rd Qu.: 0.18638
```

```
## Max. : 5.591971 Max. : 39.42090 Max. : 27.20284
## V22 V23 V24
## Min. :-10.933144 Min. :-44.80774 Min. :-2.83663
## 1st Qu. : -0.542350 1st Qu. : -0.16185 1st Qu. : -0.35459
## Median : 0.006782 Median : -0.01119 Median : 0.04098
## Mean : 0.000000 Mean : 0.00000 Mean : 0.00000
## 3rd Qu. : 0.528554 3rd Qu. : 0.14764 3rd Qu. : 0.43953
## Max. : 10.503090 Max. : 22.52841 Max. : 4.58455
## V25 V26 V27
## Min. :-10.29540 Min. :-2.60455 Min. :-22.565679
## 1st Qu. : -0.31715 1st Qu. : -0.32698 1st Qu. : -0.070840
## Median : 0.01659 Median : -0.05214 Median : 0.001342
## Mean : 0.00000 Mean : 0.00000 Mean : 0.000000
## 3rd Qu. : 0.35072 3rd Qu. : 0.24095 3rd Qu. : 0.091045
## Max. : 7.51959 Max. : 3.51735 Max. : 31.612198
## V28 Amount Class
## Min. :-15.43008 Min. : 0.00 Min. :0.000000
## 1st Qu. : -0.05296 1st Qu. : 5.60 1st Qu. :0.000000
## Median : 0.01124 Median : 22.00 Median :0.000000
## Mean : 0.00000 Mean : 88.35 Mean :0.001728
## 3rd Qu. : 0.07828 3rd Qu. : 77.17 3rd Qu. :0.000000
## Max. : 33.84781 Max. :25691.16 Max. :1.000000
```

We can see that the distribution of many PCA components is centered around zero, suggesting that the variables were standardized as part of the PCA transform. Also none of the columns have inconsistent datatypes so no conversions are required.

*# Perform basic data cleansing by checking for nulls in the dataset.*  
`colSums(is.na(df))`

```
## Time V1 V2 V3 V4 V5 V6 V7 V8 V9 V10
## 0 0 0 0 0 0 0 0 0 0 0
## V11 V12 V13 V14 V15 V16 V17 V18 V19 V20 V21
## 0 0 0 0 0 0 0 0 0 0 0
## V22 V23 V24 V25 V26 V27 V28 Amount Class
## 0 0 0 0 0 0 0 0 0 0 0
```

*# Observed no null values and hence no null treatment is required.*

There are no missing values in our dataset.

*# The following code will convert our dependent variable (Class) to a factor.*

```
df$Class = factor(df$Class)
table(df$Class)
```

```
##
## 0 1
## 284315 492
```

```
prop.table(table(df$Class))
```

```
##
## 0 1
## 0.998272514 0.001727486
```

This dataframe contains 492 fraud and 284,315 legal transactions. This shows a highly imbalanced set to work with as the frauds account for 0.172% of all transactions. We will address the problem of training a model to perform against highly imbalanced data and outline some techniques and expectations.

Imbalanced data typically refers to a problem with classification problems where the classes are not represented equally. We have a binary classification problem with 284,807. A total of 284,315 Legal labeled with Class '0' and the remaining 492 are labeled with Class '1'. This is an imbalanced dataset and the ratio Class '1' to Class '0' is  $\sim 1:578$ ! If we to use Logistic Regression, the result will ignore fraud Class and most of the time the result is overfit to the legal transaction Class. More interestingly, there are problems where a class imbalance is not just common, it is expected! For example, in datasets like this one that characterize fraudulent transaction are imbalanced. The vast majority of the transactions will be in the legal Class and a very small minority will be in the Fraud Class. It is very common to start with classification accuracy, because it is often the first measure we use when evaluating models on our classification problems. Accuracy is not appropriate here as even a classifier which labels all transactions as non-fraudulent will have over 99% accuracy. Why is that? Because our models look at the data and cleverly decide that the best thing to do is to always predict 'Legal' Class and achieve high accuracy. This is exactly why Machine Learning does not work well with imbalanced data. So now that we understand what class imbalance is and why it provides misleading classification accuracy, we will need to change our performance metric. Next, we explore our options.

## 2.2 Preparing the Data

### 2.2.1 Re-sampling the Dataset

Machine Learning algorithms assume that the dataset has balanced class distributions. As mentioned above, Machine Learning algorithms struggle with accuracy because of the unequal distribution in dependent variable. This causes the performance of existing classifiers to get biased towards majority class. We can use Sampling to build a more balanced data. Sampling technique should only be applied to the training set and not to the testing set. The three main methods are:

- **Over-sampling:** Add copies of instances from the under-represented class.
- **Under-sampling:** Delete instances from the over-represented class.
- **Synthetic Data Generation:** Randomly sample the attribute from instances in the minority Class.
  - **SMOTE** Draws artificial samples by choosing points that lie on the line connecting the rare observation to one of its nearest neighbors.
  - **ROSE** Uses smoothed bootstrapping to draw artificial samples from the neighborhood around the minority class.

## 2.3 Splitting the Data into Train and Test Sets

We will split the dataset into train and test sets in 70:30 ratio respectively. We make the decision to remove 'Time' feature from our set prior to splitting. Time feature does not indicate the actual time of the transaction and is more of listing the data in chronological order. We assume that 'Time' feature has little or no significance in correctly classifying a fraud transaction and hence eliminate this column from further analysis.

*#Remove 'Time' variable*

```
df <- df[, -1]
```

*#Change Class variable to factor*

```
df$Class <- as.factor(df$Class)
```

```
levels(df$Class) <- c("Legal", "Fraud")
```

*#Scale numeric variables*

```
df[, -30] <- scale(df[, -30])
```

```
head(df)
```

```
##           V1           V2           V3           V4           V5           V6
## 1 -0.6942411 -0.04407485  1.6727706  0.9733638 -0.245116153  0.34706734
## 2  0.6084953  0.16117564  0.1097969  0.3165224  0.043483276 -0.06181986
## 3 -0.6934992 -0.81157640  1.1694664  0.2682308 -0.364571146  1.35145121
## 4 -0.4933240 -0.11216923  1.1825144 -0.6097256 -0.007468867  0.93614819
## 5 -0.5913287  0.53154012  1.0214099  0.2846549 -0.295014918  0.07199846
```

```
## 6 -0.2174742 0.58167387 0.7525841 -0.1188331 0.305008424 -0.02231344
##          V7          V8          V9          V10          V11          V12
## 1 0.1936786 0.08263713 0.3311272 0.08338540 -0.5404061 -0.6182946
## 2 -0.0637001 0.07125336 -0.2324938 -0.15334936 1.5800001 1.0660867
## 3 0.6397745 0.20737237 -1.3786729 0.19069928 0.6118286 0.0661365
## 4 0.1920703 0.31601704 -1.2625010 -0.05046786 -0.2218912 0.1783707
## 5 0.4793014 -0.22650983 0.7443250 0.69162382 -0.8061452 0.5386257
## 6 0.3849353 0.21795429 -0.5176177 -0.34110050 1.3140441 0.3601815
##          V13          V14          V15          V16          V17          V18
## 1 -0.9960972 -0.3246096 1.6040110 -0.5368319 0.24486302 0.03076988
## 2 0.4914173 -0.1499822 0.6943592 0.5294328 -0.13516973 -0.21876220
## 3 0.7206986 -0.1731136 2.5629017 -3.2982296 1.30686559 -0.14478974
## 4 0.5101678 -0.3003600 -0.6898362 -1.2092939 -0.80544323 2.34530040
## 5 1.3522420 -1.1680315 0.1913231 -0.5152042 -0.27908030 -0.04556892
## 6 -0.3597909 -0.1430569 0.5655061 0.4584589 -0.06844494 0.08190778
##          V19          V20          V21          V22          V23          V24
## 1 0.49628116 0.32611744 -0.02492332 0.382853766 -0.17691102 0.1105067
## 2 -0.17908573 -0.08961071 -0.30737626 -0.880075209 0.16220090 -0.5611296
## 3 -2.77855597 0.68097378 0.33763110 1.063356404 1.45631719 -1.1380901
## 4 -1.51420227 -0.26985475 -0.14744304 0.007266895 -0.30477601 -1.9410237
## 5 0.98703556 0.52993786 -0.01283920 1.100009340 -0.22012301 0.2332497
## 6 -0.04077658 0.11021522 -0.28352172 -0.771425648 -0.04227277 -0.6132723
##          V25          V26          V27          V28          Amount Class
## 1 0.2465850 -0.3921697 0.33089104 -0.06378104 0.24496383 Legal
## 2 0.3206933 0.2610690 -0.02225564 0.04460744 -0.34247394 Legal
## 3 -0.6285356 -0.2884462 -0.13713661 -0.18102051 1.16068389 Legal
## 4 1.2419015 -0.4602165 0.15539593 0.18618826 0.14053401 Legal
## 5 -0.3952009 1.0416095 0.54361884 0.65181477 -0.07340321 Legal
## 6 -0.4465828 0.2196368 0.62889938 0.24563577 -0.33855582 Legal

# Split dataset into train and test sets in 70:30 ratio respectively
set.seed(42)
split <- sample.split(df$Class, SplitRatio = 0.7)
train <- subset(df, split == TRUE)
test <- subset(df, split == FALSE)
```

## 2.4 Creating Sampled Sets of Data

We will now use the sampling techniques defined above in order to produce different versions of the train set.

```
# Create Original Train Set
table(train$Class)

##
## Legal Fraud
## 199020 344

# Create Under-sampling Train Set
set.seed(42)
downsamp_train <- downSample(x = train[, -ncol(train)], y = train$Class)
table(downsamp_train$Class)

##
## Legal Fraud
## 344 344

# Create Over-sampling Train Set
```



```

set.seed(42)
upsamp_train <- upSample(x = train[, -ncol(train)], y = train$Class)
table(upsamp_train$Class)

##
## Legal Fraud
## 199020 199020

# Create SMOTE Train Set
set.seed(42)
smote_train <- SMOTE(Class ~ ., data = train)
table(smote_train$Class)

##
## Legal Fraud
## 1376 1032

# Create ROSE Train Set
set.seed(42)
rose_train <- ROSE(Class ~ ., data = train)$data
table(rose_train$Class)

##
## Legal Fraud
## 99791 99573

```

The Following table summarizes the different count of Class variable (Legal/Fraud) in our train set.

Sampling Methods (Train set)	# of Legal	# of Fraud
No Sampling: Original	199020	344
Under Sampling	344	344
Over Sampling	199020	199020
SMOTE	1376	1032
ROSE	99791	99573

### 3 Modeling Approach

As said before, Accuracy is not the metric to use when working with an imbalanced dataset. There are metrics that have been designed to tell a better story when working with imbalanced classes. ROC Curves and Precision-Recall Curves are some useful tools. ROC curves are appropriate when the observation are balanced between each class, whereas precision-recall curves are appropriate for imbalanced datasets. Since we have created a balanced data using the sampling technique above, we will continue with ROC curves for predicting the probability of a binary outcome. ROC curve stands for Receiver Operating Characteristic curve. The plot depicts the false positive rate on x-axis vs the true positive rate (y-axis) for a number of different candidate threshold (between 0 and 1). In other words, it plots the false alarm rate versus the hit rate. We will explore ROC Curves AUC (Area Under Curve). The higher the AUC, the better the performance of the model at distinguishing between the positive and negative classes. The AUC value lies between 0.5 and 1 where 0.5 denotes a bad classifier and 1 denotes an excellent classifier. Our modeling approach will involve training a single classifier on the train set with class imbalance suitably altered using each of the train set versions above. Depending on which version of train set yields the best roc-auc score on a holdout test set. We will then build subsequent models using that chosen technique. In our analysis, we will train several algorithms such as CART, GLM, Random Forest, and XGBoost.


### 3.1 CART Method

We choose CART (Classification And Regression Tree) as first model. Before we start using sampling let us first look at how CART performs with imbalanced data. we use the function `roc.curve` available in the ROSE package to gauge model performance on the test set.

#### 3.1.1 CART Method: Calculate AUC using Original Train Dataset

```
# CART Model Performance on original imbalanced dataset
set.seed(42)
original_fit <- rpart(Class ~ ., data = train)
#Evaluate Model Performance on test set
pred_original <- predict(original_fit, newdata = test, method = "class")
ROSE::roc.curve(test$Class, pred_original[,2], plotit = TRUE)
```

ROC curve



```
## Area under the curve (AUC): 0.898
```

We evaluate the model performance on test data by calculating the roc auc score. AUC score on the original dataset is **0.898**. This is our first attempt at our AUC and we believe we can improve on this score.


Method	AUC
CART Method: Original	<b>0.898</b>

#### 3.1.2 CART Method: Calculate AUC using Under-sampled Train Dataset

```
set.seed(42)
# Build down-sampled model
downsample_fit <- rpart(Class ~ ., data = downsamp_train)
predict_down <- predict(downsample_fit, newdata = test)
print("Fitting downsampled model to test data")
```

```
## "Fitting downsampled model to test data"
ROSE::roc.curve(test$Class, predict_down[,2], plotit = TRUE)
```

**ROC curve**




```
## Area under the curve (AUC): 0.913
```

Method	AUC
CART Method: Original	0.898
CART Method: Undersampling	<b>0.913</b>

### 3.1.3 CART Method: Calculate AUC using Over-sampled Train Dataset

```
set.seed(42)
# Build up-sampled model
upsamp_fit <- rpart(Class ~ ., data = upsamp_train)
predict_up <- predict(upsamp_fit, newdata = test)
print("Fitting upsampled model to test data")
## "Fitting upsampled model to test data"
ROSE::roc.curve(test$Class, predict_up[,2], plotit = TRUE)
```



## Area under the curve (AUC): 0.935


Method	AUC
CART Method: Original	0.898
CART Method: Undersampling	0.913
CART Method: Oversampling	<b>0.935</b>

### 3.1.4 CART Method: Calculate AUC using SMOTE sampled Train Dataset

```
set.seed(42)
# train the models
smote_fit <- rpart(Class ~ ., data = smote_train)
pred_smote <- predict(smote_fit, newdata = test)
print("Fitting smote model to test data")

## "Fitting smote model to test data"

#try and predict an outcome from the test set
ROSE::roc.curve(test$Class, pred_smote[,2], plotit = TRUE)
```



## Area under the curve (AUC): 0.908


Method	AUC
CART Method: Original	0.898
CART Method: Undersampling	0.913
CART Method: Oversampling	0.935
CART Method: SMOTE	<b>0.908</b>

### 3.1.5 CART Method: Calculate AUC using ROSE sampled Train Dataset

```
set.seed(42)
# # train the models. Build rose model
rose_fit <- rpart(Class ~ ., data = rose_train)
pred_rose <- predict(rose_fit, newdata = test)
print("Fitting rose model to test data")

## "Fitting rose model to test data"

# try and predict an outcome from the test set
ROSE::roc.curve(test$Class, pred_rose[,2], plotit = TRUE)
```




## Area under the curve (AUC): 0.919

Method	AUC
CART Method: Original	0.898
CART Method: Undersampling	0.913
CART Method: Oversampling	0.935
CART Method: SMOTE	0.908
CART Method: ROSE	<b>0.919</b>

We see that all the sampling techniques have yielded better AUC scores than the simple imbalanced dataset. We will test different models now using the up sampling technique as that has given the highest AUC score.

### 3.2 Logistic Regression (GLM) Model

```
# train the models
fit_glm <- glm(Class ~ ., data = upsamp_train, family = "binomial")
predict_glm <- predict(fit_glm, newdata = test, type = "response")
# try and predict an outcome from the test set
ROSE::roc.curve(test$Class, predict_glm, plotit = TRUE)
```




## Area under the curve (AUC): 0.967

Method	AUC
CART Method: Original	0.898
CART Method: Undersampling	0.913
CART Method: Oversampling	0.935
CART Method: SMOTE	0.908
CART Method: ROSE	0.919
GLM Fit	<b>0.967</b>

### 3.3 Random Forest (RF Fit) Model

```
# train the models
x = upsamp_train[,-30]
y = upsamp_train[,30]
fit_rf <- Rborist(x, y, ntree = 1000, minNode = 21, maxLeaf = 12)
predict_rf <- predict(fit_rf, test[,-30], ctgCensus = "prob")
prob <- predict_rf$prob
# try and predict an outcome from the test set
ROSE::roc.curve(test$class, prob[,2], plotit = TRUE)
```



## Area under the curve (AUC): 0.971

Method	AUC
CART Method: Original	0.898
CART Method: Undersampling	0.913
CART Method: Oversampling	0.935
CART Method: SMOTE	0.908
CART Method: ROSE	0.919
GLM Fit	0.967
RF Fit	<b>0.971</b>

### 3.4 xgboost (XGB Fit) Model

Lastly, we can also try XGBoost, which is based on Gradient Boosted Trees and is a more powerful model compared to both Logistic Regression and Random Forest. Now that we have seen how to evaluate models on this dataset, let's look at how we can use a final model to make predictions.

```
# train the models
labels <- upsamp_train$Class
y <- recode(labels, "Legal" = 0, "Fraud" = 1)
xgb <- xgboost(data = data.matrix(upsamp_train[, -30]),
  label = y,
  eta = 0.1,
  gamma = 0.1,
  max_depth = 10,
  nrounds = 400,
  objective = "binary:logistic",
  colsample_bytree = 0.5,
  verbose = 0,
```




```

nthread = 8,
seed = 42
)

## [21:01:40] WARNING: amalgamation/./src/learner.cc:1095: Starting in XGBoost 1.3.0, the default eval
predict_xgb <- predict(xgb, data.matrix(test[, -30]))
#try and predict an outcome from the test set
ROSE::roc.curve(test$Class, predict_xgb, plotit = TRUE)

```

**ROC curve**



```
## Area under the curve (AUC): 0.982
```

Method	AUC
CART Method: Original	0.898
CART Method: Undersampling	0.913
CART Method: Oversampling	0.935
CART Method: SMOTE	0.908
CART Method: ROSE	0.919
GLM Fit	0.967
RF Fit	0.971
XGB Fit	<b>0.982</b>

XGB can also automatically provide estimates of feature importance from a trained predictive model. Feature importance scores can provide insight into the dataset. This means we could save disk space and computation time by only training the model on the most correlated/important variables.

## 4 Results

We see that all the sampling techniques have yielded better AUC scores than the simple imbalanced dataset. We then tested different models using the up sampling technique as that has given the highest AUC score. With an auc score of 0.982 the XGBOOST model has performed the best though both the random forest and logistic regression models have shown reasonable performance. This also indicated that ROC Curve AUC has been a good performance metric as it allow us to compare variety of Machine Learning algorithms to achieve AUC closer to 1. It has been shown that even a very simple logistic regression model can achieve good result, while a much more complex Random Forest model improves upon logistic regression in terms of AUC. However, XGBoost model improves upon both models.

Method	AUC
CART Method: Original	0.898
CART Method: Undersampling	0.913
CART Method: Oversampling	0.935
CART Method: SMOTE	0.908
CART Method: ROSE	0.919
GLM Fit	0.967
RF Fit	0.971
XGB Fit	<b>0.982</b>

## 5 Conclusion

This project has explored the task of identifying fraudulent transactions based on a dataset of anonymized features. We have studied sampling as a way to deal with unbalanced datasets and discussed why accuracy is not an appropriate measure of performance and we used the metric AUC ROC to evaluate how sampling can lead to better training. We concluded that the oversampling technique works best on the dataset and has achieved significant improvement in model performance over the imbalanced data. The best score of 0.982 was obtained by XGBOOST model. RF and Logistic models also performed fairly well. using an XGBOOST model though both random forest and logistic regression models performed well too. This model has many valid real-world use cases. For example, a bank could take a similar approach and reduce the amount of money spent trying to detect fraud by automating it with a machine. We can protect consumers with this technology by having a similar model integrated into the transaction process to notify the consumer and

bank of fraud within minutes, rather than days. This project has explored the task of identifying fraudulent transactions based on a dataset of anonymized features.

A future task to explore is to focus on two metrics: ROC Curves and Precision-Recall Curves to use on an imbalanced dataset and compare the AUC score. Learn more about, does it make sense to create balance set or use tools like Precision Recall on an imbalanced data.